

ask.py

license MIT

A single Python program to implement the search-extract-summarize flow, similar to AI search engines such as Perplexity.

- You can run it on command line or with a GradIO UI.
- You can control the output behavior, e.g., extract structured data or change output language,
- You can control the search behavior, e.g., restrict to a specific site or date, or just scrape a specified list of URLs.
- You can run it in a cron job or bash script to automate complex search/data extraction tasks.

We have a running UI example [in HuggingFace Spaces](#).

Ask.py - Web Search-Extract-Summarize

Search the web with the query and summarize the results. Source code: <https://github.com/pengfeng/ask.py>

Query

What is RAG?

Date Restrict (Optional) [0 or empty means no date limit.]

0

Target Sites (Optional) [Empty means search the whole web.]

Output Language (Optional) [Default is English.]

English

Output Length in words (Optional) [Default is automatically decided by LLM.]

0

URL List (Optional) [When specified, scrape the urls instead of searching the web.]

Additional Inputs

Clear

Submit

output

RAG, or Retrieval-Augmented Generation, is an AI framework that combines traditional information retrieval systems with generative large language models (LLMs) to enhance the accuracy and relevance of responses generated for knowledge-intensive tasks. It operates through a two-step process: first, it retrieves relevant information from external data sources, and then it uses that information to generate coherent and contextually appropriate responses. This method allows for more accurate and up-to-date outputs as it integrates real-time knowledge into the generative process, effectively addressing the limitations of static training data in LLMs[3][5][10].

References

[1] <https://stackoverflow.blog/2023/10/18/retrieval-augmented-generation-keeping-llms-relevant-and-current/>
[2] <https://stackoverflow.blog/2023/10/18/retrieval-augmented-generation-keeping-llms-relevant-and-current/>
[3] <https://stackoverflow.blog/2023/10/18/retrieval-augmented-generation-keeping-llms-relevant-and-current/>
[4] <https://cloud.google.com/use-cases/retrieval-augmented-generation>
[5] <https://cloud.google.com/use-cases/retrieval-augmented-generation>
[6] <https://www.cohesity.com/glossary/retrieval-augmented-generation-rag/>
[7] <https://stackoverflow.blog/2023/10/18/retrieval-augmented-generation-keeping-llms-relevant-and-current/>
[8] <https://www.cohesity.com/glossary/retrieval-augmented-generation-rag/>
[9] <https://stackoverflow.blog/2023/10/18/retrieval-augmented-generation-keeping-llms-relevant-and-current/>
[10] <https://www.cohesity.com/glossary/retrieval-augmented-generation-rag/>

Report Error

[!NOTE]

- Our main goal is to illustrate the basic concepts of AI search engines with the raw constructs.
Performance or scalability is not in the scope of this program.
- We are planning to open source a real search-enabled AI toolset with real DB setup, real document pipeline, and real query engine soon. Star and watch this repo for updates!

[UPDATE]

- 2024-11-10: add Chonkie as the default chunker
- 2024-10-28: add extract function as a new output mode
- 2024-10-25: add hybrid search demo using DuckDB full-text search
- 2024-10-22: add GradIO integration
- 2024-10-21: use DuckDB for the vector search and use API for embedding
- 2024-10-20: allow to specify a list of input urls
- 2024-10-18: output-language and output-length parameters for LLM
- 2024-10-18: date-restrict and target-site parameters for search

The search-extract-summarize flow

Given a query, the program will

- search Google for the top 10 web pages
- crawl and scrape the pages for their text content
- chunk the text content into chunks and save them into a vectordb
- perform a vector search with the query and find the top 10 matched chunks
- [Optional] search using full-text search and combine the results with the vector search
- [Optional] use a reranker to re-rank the top chunks
- use the top chunks as the context to ask an LLM to generate the answer
- output the answer with the references

Of course this flow is a very simplified version of the real AI search engines, but it is a good starting point to understand the basic concepts.

One benefit is that we can manipulate the search function and output format.

For example, we can:

- search with date-restrict to only retrieve the latest information.
- search within a target-site to only create the answer from the contents from it.
- ask LLM to use a specific language to answer the question.
- ask LLM to answer with a specific length.
- crawl a specific list of urls and answer based on those contents only.

Quick start

```
# recommend to use Python 3.10 or later and use venv or conda to create a virtual
environment
% pip install -r requirements.txt

# modify .env file to set the API keys or export them as environment variables as
below

# right now we use Google search API
% export SEARCH_API_KEY="your-google-search-api-key"
```

```
% export SEARCH_PROJECT_KEY="your-google-cx-key"

# right now we use OpenAI API
% export LLM_API_KEY="your-openai-api-key"

# By default, the program will start a web UI. See Gradio Deployment section for
more info.

# Run the program on command line with -c option
% python ask.py -c -q "What is an LLM agent?"

# we can specify more parameters to control the behavior such as date_restrict and
target_site
% python ask.py --help
Usage: ask.py [OPTIONS]

    Search web for the query and summarize the results.

Options:
  -q, --query TEXT                Query to search
  -o, --output-mode [answer|extract]
                                   Output mode for the answer, default is a
                                   simple answer
  -d, --date-restrict INTEGER      Restrict search results to a specific date
                                   range, default is no restriction
  -s, --target-site TEXT           Restrict search results to a specific site,
                                   default is no restriction
  --output-language TEXT           Output language for the answer
  --output-length INTEGER          Output length for the answer
  --url-list-file TEXT             Instead of doing web search, scrape the
                                   target URL list and answer the query based
                                   on the content
  --extract-schema-file TEXT       Pydantic schema for the extract mode
  --inference-model-name TEXT      Model name to use for inference
  --hybrid-search                  Use hybrid search mode with both vector
                                   search and full-text search
  -c, --run-cli                   Run as a command line tool instead of
                                   launching the Gradio UI
  -l, --log-level [DEBUG|INFO|WARNING|ERROR]
                                   Set the logging level [default: INFO]
  --help                          Show this message and exit.
```

Libraries and APIs used

- [Google Search API](#)
- [OpenAI API](#)

- [Jinja2](#)
- [bs4](#)
- [DuckDB](#)
- [GradIO](#)
- [Chonkie](#)

GradIO Deployment

[!NOTE]

Original GradIO app-sharing document [here](#).

Quick test and sharing

By default, the program will start a web UI and share through GradIO.

```
% python ask.py
* Running on local URL:  http://127.0.0.1:7860
* Running on public URL: https://77c277af0330326587.gradio.live

# you can also specify SHARE_GRADIO_UI to only run locally
% export SHARE_GRADIO_UI=False
% python ask.py
* Running on local URL:  http://127.0.0.1:7860
```

To share a more permanent link using HuggingFace Spaces

- First, you need to [create a free HuggingFace account](#).
- Then in your [settings/token page](#), create a new token with Write permissions.
- In your terminal, run the following commands in you app directory to deploy your program to HuggingFace Spaces:

```
% pip install gradio
% gradio deploy
Creating new Spaces Repo in '/home/you/ask.py'. Collecting metadata, press Enter
to accept default value.
Enter Spaces app title [ask.py]: ask.py
Enter Gradio app file [ask.py]:
Enter Spaces hardware (cpu-basic, cpu-upgrade, t4-small, t4-medium, 14x1, 14x4,
zero-a10g, a10g-small, a10g-large, a10g-largex2, a10g-largex4, a100-large, v5e-
1x1, v5e-2x2, v5e-2x4) [cpu-basic]:
Any Spaces secrets (y/n) [n]: y
Enter secret name (leave blank to end): SEARCH_API_KEY
Enter secret value for SEARCH_API_KEY: YOUR_SEARCH_API_KEY
Enter secret name (leave blank to end): SEARCH_PROJECT_KEY
```

```
Enter secret value for SEARCH_API_KEY: YOUR_SEARCH_PROJECT_KEY
Enter secret name (leave blank to end): LLM_API_KEY
Enter secret value for LLM_API_KEY: YOUR_LLM_API_KEY
Enter secret name (leave blank to end):
Create Github Action to automatically update Space on 'git push'? [n]: n
Space available at https://huggingface.co/spaces/your_user_name/ask.py
```

Now you can use the HuggingFace space app to run your queries.

Use Cases

- [Search like Perplexity](#)
- [Only use the latest information from a specific site](#)
- [Extract information from web search results](#)